

SPDB Annotation Codebook v1.1

Spanish Political Discourse Benchmark (SPDB)

Version: codebook-v1.1.0 (aligned with guidelines-v1.1)

Date: 2026-09-29

Audience: Primary annotators, lead annotators, adjudicators

Canonical ontology: labels/*.tsv (eight PF_* classes unchanged)

Supersedes for new coding: codebook-v1.0.0 / SPDB_Codebook_v1.md

Wave-1 note: Wave-1 annotations remain **v1.0-era diagnostic annotations** and are **not** reclassified under v1.1. Do not mix v1.0 and v1.1 labels in one reliability statistic.

Document control

Field	Value
Guideline version	guidelines-v1.1
Prior instrument	codebook-v1.0.0 / guidelines-v1.0 (Wave-1 diagnostic)
Label inventory	8 pragmatic functions · 7 fallacies + explicit none · 2 experimental arms (unchanged)
Unit of analysis	Discourse unit (one dominant pragmatic function per span)
Language of text	Spanish (es)
Language of label IDs	English codes (PF_*, FAL_*, ...)
Attestation	Required before reliability coding — see annotation/attestation/annotator_session_attestation.md

Reliability target (codebook design goal): Cohen's κ and Krippendorff's $\alpha > 0.80$ on pragmatic function after calibration (excellent tier per annotation/pilot_001/pilot_metrics.md). Release MVP floor remains $\kappa \geq 0.67$ and $\alpha \geq 0.67$ for the next independent reliability wave under this instrument.

1. Purpose and scope

This codebook operationalises human annotation for SPDB v1. Annotators assign:

1. **Pragmatic function** — exactly one PF_* label (required, primary supervision).
2. **Political fallacies** — zero to three FAL_* labels, or explicit none (required).
3. **Experimental labels** — SV_* and CA_* when experimental_subset=true (optional arm).

Labels apply to the **entire discourse unit** in the task, not to a sentence fragment inside it. Metadata (speaker_name, speaker_party, date, source_type) provides context only.

Independence rule: Pragmatic function and fallacy labels are coded **independently**. A evidence-based PF_ATTACK may co-occur with no fallacy; PF_ATTACK may co-occur with FAL_ADHOM.

2. Unit of analysis

A **discourse unit** is the smallest contiguous span an informed reader treats as carrying **one dominant pragmatic function** toward a political target (actor, policy, institution, or electorate).

Rule	Specification
Span	Full text field in Label Studio / CSV row
Overlap	Units do not overlap; adjacent units do not share sentences
Length	20–2,000 characters after normalisation (pre-annotation pool)
External context	Not allowed (no web search, no speaker biography beyond supplied metadata, no political-world knowledge used to invent unstated prior criticism)

2.1 Context policy (v1.1)

Annotators **may** use:

1. The unit text (required).
2. Supplied metadata fields: speaker_name, speaker_party, date, document_id / session identifiers shown in the task — as **orientation only**.
3. **±1 adjacent discourse unit** from the **same** document_id when displayed in the interface (or provided in the worksheet), for:
 - fallacy coding;
 - deciding PF_DEFENSE vs PF_ATTACK / PF_ADVOCACY / PF_INFO when reactivity must be recovered;
 - other borderline PF decisions under §4.1.
4. The **pending question / agenda item** if it is visible in the ±1 window or in the unit itself (e.g. chair naming the punto del orden del día).

Annotators **must not**:

1. Let speaker_name alone determine the PF label (including Presidencia). Labels follow **utterance function**, not office identity (§4.2).
2. Use party identity, ideology, or reputation as a shortcut for PF.
3. Search the web, Hansard beyond the provided ±1 window, news, or social media.
4. Infer a prior attack that is **not** recoverable from the unit + allowed ±1 context.

If allowed context is still insufficient to choose between two PF labels: select the best label under the decision tree, enable **Borderline PF case** (borderline=true), and continue. Do **not** invent a ninth output label. Do **not** leave PF empty.

3. Annotation workflow (summary)

Primary coding (blind) → disagreement detection → lead third pass → PI adjudication → gold label

Step	Actor	Action
1	Annotator A & B	Independent primary labels on double-coded sample (20%) or single pass
2	Pipeline	Flag PF or fallacy mismatches; flag borderline=true
3	Lead annotator	Third pass on disagreements
4	PI / adjudicator	Gold label within 72 h if disagreement persists
5	Export	adjudicated=true; log in adjudication_log.jsonl (IDs only)

Full UI order: Context → (Adjudication panel if applicable) → PF → Fallacies → Experimental → Quality flags.

4. Master decision tree — pragmatic function

Apply **top to bottom**; assign the **first** category whose dominant-move criterion is satisfied. If two categories tie, use the tie-breakers in §4.1 and pair rules in §4.2–§4.3.

START: Read full unit (+ allowed context, §2.1). Ignore speaker biography beyond metadata.

1. Is the dominant move CHAMBER MANAGEMENT of the proceeding itself (order of business, turn-taking, voting instructions/results, timekeeping, suspension/resumption, agenda transition, rule reminders, recognition of speakers) – i.e. managing how the sitting runs – rather than informing about bill/policy substance or arguing a political position?
YES → PF_PROCEDURAL
NO → continue
(See §4.2. Speaker identity alone never decides this step.)
2. Does the speaker avoid the pending question/topic and pivot to a different issue without substantive engagement?
YES → PF_DEFLECT
NO → continue
3. Does the span commit to a future legislative/government action by this actor (presentaremos, nos comprometemos, votaremos a favor/en contra as pledge)?
YES → PF_PROPOSAL
NO → continue
4. Is the dominant move mobilising citizens/voters via hope, fear, solidarity, or shared identity (not primarily policy detail or opponent attack)?
YES → PF_APPEAL
NO → continue
5. Is the speaker primarily rebutting criticism directed at self/party/policy (reactive defense), with the prior criticism recoverable from the unit or allowed ±1 context (§2.1)?
YES → PF_DEFENSE
NO → continue

6. Is the primary target a political actor/group's competence, integrity, record, or legitimacy (not merely abstract policy)?
YES → PF_ATTACK
NO → continue
7. Does the span argue for/against a policy/programme/normative goal as the dominant move?
YES → PF_ADVOCACY
NO → continue
8. Does the span convey facts, figures, or bill/policy content with minimal evaluative framing as the dominant move (NOT chamber-management reporting)?
YES → PF_INFO
NO → re-read; apply §4.1–§4.3; flag borderline if still tied

4.1 Tie-breakers (PF)

1. Choose the function directed at the **primary political target** in the span.
2. If courtesy opener + substance: label the **substance**, not the opener.
3. If attack and advocacy coexist: whichever occupies **more rhetorical labour** (sentences / clauses devoted to the move). Count roughly; do not obsess over exact tokens.
4. If attack and appeal coexist: choose **APPEAL** only if mobilisation of citizens/voters is the dominant move; otherwise **ATTACK** (see §4.3.2).
5. If still tied: select best label and enable **Borderline PF case** (borderline=true); expected ≤ **5%** of units.
6. Never assign two PF labels; never leave PF empty.

4.2 Operational rule — PF_PROCEDURAL vs PF_INFO (mandatory)

Principle: Classify by the **function of the utterance**, not by who speaks.

Code as PF_PROCEDURAL when the dominant move is...	Examples of moves
Managing how the sitting runs	Order-of-business announcements; agenda transitions; naming the punto del orden del día to open/close it
Turn-taking / recognition	“Tiene la palabra...”, assigning turns, closing a turn
Voting management	Opening/closing a vote; reporting vote tallies as chamber result ; “queda aprobado/rechazado” as procedural outcome
Timekeeping / discipline	Speaking-time limits; calls to order; photography bans; seating instructions
Suspension / resumption	Suspending, resuming, or closing the sitting
Formal chamber management	Reminders of Rules of Procedure that organise the debate without arguing policy merits
Factual statements about the proceeding itself	“Efectuada la votación, dio el siguiente resultado: ...”; announcing which group speaks next

Code as PF_INFO when the dominant move is...	Examples of moves
Reporting bill / policy / subject-matter content	Summarising articles of a proposition; citing socio-economic figures; describing programme measures

Code as PF_INF0 when the dominant move is...	Examples of moves
Technical exposition with minimal evaluation	Process descriptions of a policy domain (not chamber ritual)

Chair / Presidencia turns: Being chair raises the **prior probability** of PF_PROCEDURAL but is **not decisive**. If a chair (or any speaker) utters **substantive policy/bill information** or political argument as the dominant move, code that function (PF_INF0, PF_ADVOCACY, etc.), not PF_PROCEDURAL.

Mixed turns (procedural shell + substance):

1. Apply §4.1.2: if a short procedural formula opens a longer substantive speech, label the **substance**.
2. If the unit is **only** chamber management (including vote results / speaker recognition), label PF_PROCEDURAL even when it contains numbers (tallies are proceeding-facts, not policy-facts).
3. **Tie-break:** if $\geq 50\%$ of the rhetorical labour manages the sitting $\rightarrow$ PF_PROCEDURAL; if $\geq 50\%$ reports policy/bill substance with minimal evaluation $\rightarrow$ PF_INF0; if still tied $\rightarrow$ `borderline=true` and prefer PF_PROCEDURAL when the labour is sitting-management of the **current** punto, else PF_INF0.

Terminology ban (v1.1): Do **not** treat “procedural information” as a synonym of PF_INF0. Chamber-procedure reporting belongs under PF_PROCEDURAL. PF_INF0 is reserved for **subject-matter / bill / factual policy** information.

4.3 High-value pair discriminators

4.3.1 PF_ATTACK vs PF_ADVOCACY

	ATTACK	ADVOCACY
Discriminating question	Is the primary target an actor/group (competence, integrity, record, legitimacy)?	Is the primary target a policy/programme/normative design ?
Positive criterion	Charges, record comparisons, legitimacy challenges aimed at who is unfit/wrong	Args for/against what should be done, including policy criticism without personal targeting
Common trap	Treating any criticism of a government bill as attack, or any “ustedes” as automatic attack	
Tie-break	More rhetorical labour on actor vs policy (§4.1.3); if equal $\rightarrow$ <code>borderline=true</code>	
Context	Usually unit-local; ± 1 only if referent of “ustedes/ellos” is otherwise opaque	

4.3.2 PF_ATTACK vs PF_APPEAL

	ATTACK	APPEAL
Discriminating question	Does the span mainly delegitimise opponents ?	Does it mainly mobilise citizens/voters via hope, fear, solidarity, identity?

	ATTACK	APPEAL
Positive criterion	Opponent-directed competence/integrity/record	Audience-directed mobilisation (“españoles”, “juntos”, civic duty)
Common trap	Hope + blame blends in campaign rhetoric	
Tie-break	If mobilisation of the electorate is dominant → APPEAL; if opponent targeting is dominant → ATTACK; equal → <code>borderline=true</code>	
Context	Unit-local usually sufficient	

4.3.3 PF_ADVOCACY VS PF_PROPOSAL

	ADVOCACY	PROPOSAL
Discriminating question	Does the speaker argue a position without pledging an action by this actor?	Does the speaker commit this actor to a future action?
Positive criterion	Should / must / we support / we oppose (without pledge)	presentaremos, nos comprometemos, impulsaremos, anunciamos que presentaremos...
Common trap	Soft future talk (“habría que”) coded as proposal; or voting “a favor” as position without new pledge	
Tie-break	Explicit commitment verb/pledge by this actor → PROPOSAL; else ADVOCACY. Chamber “vamos a votar” from the chair is PROCEDURAL (§4.2), not PROPOSAL	
Context	Unit-local	

4.3.4 PF_ATTACK VS PF_DEFENSE

	ATTACK	DEFENSE
Discriminating question	Is the speaker initiating criticism of another actor?	Is the speaker reacting to criticism of self/party/policy?
Positive criterion	Offensive charge toward others	Denial, correction, shielding of own side after recoverable prior criticism
Common trap	Coding counter-charge as defense when no prior criticism is recoverable	
Tie-break	If prior criticism is recoverable in unit/±1 → DEFENSE when rebuttal dominates; if the span mainly launches new charges → ATTACK	
Context	±1 required when reactivity is unclear from the unit alone (§2.1). If still unrecoverable → do not force DEFENSE; prefer ATTACK/ADVOCACY/INFO as tree indicates and set <code>borderline=true</code> if tied	

4.3.5 PF_DEFLECT VS PF_DEFENSE

	DEFLECT	DEFENSE
Discriminating question	Does the speaker leave the pending issue unanswered and change subject?	Does the speaker address the challenge (deny/correct/shield)?
Positive criterion	Pivot / whatabout / side-story without engagement	On-topic rebuttal
Common trap	Treating any counter-attack as deflection	
Tie-break	On-topic engagement → not DEFLECT (use DEFENSE or ATTACK). Unanswered pending question + topic shift → DEFLECT	
Context	Pending question from ±1 / agenda metadata when needed. DEFLECT is rare in prepared plenary speech	

4.3.6 PF_APPEAL VS PF_ADVOCACY / PF_PROPOSAL

	APPEAL	ADVOCACY / PROPOSAL
Discriminating question	Is the dominant move mobilising citizens/ voters ?	Is it policy argument (advocacy) or action pledge (proposal)?
Positive criterion	Hope/fear/solidarity/identity addressed to the electorate	Policy design merits, or explicit commitment by this actor
Common trap	Civic-duty slogans that also endorse a bill	
Tie-break	Electorate mobilisation dominates → APPEAL; else apply §4.3.3 for advocacy vs proposal	
Context	Unit-local	

4.3.7 PF_ATTACK VS PF_PROPOSAL / PF_INFO

	ATTACK	PROPOSAL / INFO
Discriminating question	Is the primary move delegitimising an actor ?	Is it a pledge (proposal) or subject-matter report (info)?
Positive criterion	Actor competence/integrity/record	Commitment verbs → PROPOSAL; bill/figures with minimal evaluation → INFO
Common trap	Dual-move “we will fix what they ruined” spans	
Tie-break	More labour on actor charge → ATTACK; on pledge → PROPOSAL; on description → INFO; tied → <code>borderline=true</code>	
Context	Unit-local	

5. Pragmatic function labels

Enter **label ID** only (e.g. PF_ADVOCACY), not the English name.

5.1 PF_ADVOCACY — Policy advocacy

Definition: Supports or advances a policy position, programme, or normative goal.

Decision rule: Dominant move argues **for or against a policy design**, law, or normative outcome. Criticism of **policy** (not primarily personal character) belongs here unless the span is reactive defense (PF_DEFENSE).

Example (positive):

«En Ciudadanos abogamos [...] que se tenía que definir la figura del consumidor vulnerable [...] que se tenía que fijar un criterio de renta [...]»

— Melisa Rodríguez Hernández (Cs). Unit spdb-v1-unassigned-ceada2c252a1.

Counterexample:

«Tienen ustedes la osadía de hablar de la contratación temporal cuando precisamente la reforma del año 1984 es la que rompe el principio de causalidad [...]»

— Carolina España Reina (PP). **Not advocacy** → PF_ATTACK (opponent record).

Confused with	Resolution
PF_ATTACK	See §4.3.1 — advocacy → policy ; attack → actor
PF_PROPOSAL	See §4.3.3 — proposal commits ; advocacy argues
PF_INFO	Advocacy embeds facts in normative argument; info is descriptive subject-matter content

Additional contrastive note (v1.1): Mixed actor+policy spans use rhetorical-labour tie-break (§4.1.3 / §4.3.1). Prefer `borderline=true` rather than inventing a dual label.

5.2 PF_ATTACK — Position attack

Definition: Challenges an actor's competence, integrity, legitimacy, or record.

Decision rule: Primary target is a **political actor or group**. Personalised charges, record comparisons, and legitimacy challenges qualify. Evidence-based policy disagreement **without** personal targeting → PF_ADVOCACY.

Example (positive):

«No es así, señora Arrimadas. Ustedes ni siquiera han estado sentados en la mesa donde se tomaban estas decisiones [...] tampoco han conseguido nada.»

— Pablo Echenique Robba (UP). Unit spdb-v1-unassigned-b2eaf6b9bc1c.

Counterexample:

«son tres principalmente las reformas que plantea esta proposición de ley: por una parte, permitir la rectificación registral [...]»

— Joseba Agirretxea Urresti (EAJPNV). **Not attack** → PF_INFO.

Confused with	Resolution
PF_ADVOCACY	See §4.3.1
PF_APPEAL	See §4.3.2
PF_DEFENSE	See §4.3.4 — defense reacts ; attack initiates
PF_DEFLECT	See §4.3.5 — attack engages ; deflect changes subject

5.3 PF_DEFENSE — Defense / rebuttal

Definition: Rebuts criticism or protects actor or policy reputation.

Decision rule: Speaker responds to **prior criticism** (explicit or clearly recoverable from the unit or allowed ± 1 context, §2.1) by denying, correcting, or shielding their party, government, or policy. Do **not** infer unrecovered prior attacks from political-world knowledge.

Example (positive):

«Falso, no llegan ni al 1 %.» (*rebutting opponents' fiscal claims in debate context*)

— Carolina España Reina (PP). Unit spdb-v1-unassigned-c6349a3fb522 (pre-v1.1 codebook example; retained).

Illustrative constructed example (not corpus data):

«No es verdad que hayamos recortado la sanidad; el presupuesto de este año aumenta un 4 %.» — *constructed rebuttal after an explicit prior charge in ± 1 context* → PF_DEFENSE.

Counterexample:

«nos vemos obligados a abstenernos en esta votación.»

— Íñigo Barandiaran Benito (EAJPNV). **Not defense** → position-taking / PF_ADVOCACY.

If prior criticism is not recoverable: do not default to DEFENSE; continue the tree (§4 steps 6–8) and use `borderline=true` if needed (§4.3.4).

5.4 PF_PROPOSAL — Proposal / commitment

Definition: Commits to future action, concrete measure, or implementation pledge.

Decision rule: Explicit **forward-looking commitment** (legislative initiative, amendment pledge, implementation promise). General policy praise without commitment → PF_ADVOCACY.

Example (positive):

«Desde ahora anunciamos que presentaremos una enmienda de eliminación [...] nos comprometemos a que en breve en esta Cámara se debata su toma en consideración.»

— María Galovart Carrera (PSOE). Unit spdb-v1-unassigned-46bb18b3c74a.

Counterexample:

«Comienza la votación.» — Presidencia. **Not proposal** → PF_PROCEDURAL.

5.5 PF_APEAL — Mobilization appeal

Definition: Seeks electoral or civic support via identity, values, fear, hope, or solidarity.

Decision rule: Dominant move **mobilises** audience (citizens, voters) through shared identity, hope, fear, or solidarity—not policy detail or opponent attack.

Example (positive):

«A todos los españoles queremos transmitirles un mensaje de esperanza [...] juntos, con responsabilidad, pero también con determinación.»

— María Montero Cuadrado (PSOE). Unit spdb-v1-unassigned-2063abcb1366.

Counterexample:

«Creo que ha quedado muy claro quién quiere que este país pueda salir de la crisis [...] y quién prefiere que no.»

— Same unit, different excerpt. **Not appeal** → implicit PF_ATTACK.

5.6 PF_INFO — Informational

Definition: Conveys **subject-matter** factual information (policy, bill content, figures, technical description) with minimal evaluative framing.

Decision rule: Reports facts, figures, bill contents, or technical **policy-domain** information without dominant persuasion, attack, mobilisation, or **chamber management**.

Not for order-of-business reporting, vote tallies, or speaker recognition — those are PF_PROCEDURAL (§4.2).

Example (positive) — retained from codebook-v1.0.0 (illustrative of bill exposition; not Wave-1 adjudicated gold):

«son tres principalmente las reformas que plantea esta proposición de ley [...] permitir la rectificación registral [...] suprimir la obligación de aportar documento médico [...]»

— Joseba Agirretxea Urresti (EAJPNV). Unit spdb-v1-unassigned-eb453642ca5f.

Illustrative constructed example (not corpus data):

«El artículo 3 de la proposición fija un umbral de renta de 21.000 euros y un plazo de tres meses para la solicitud.» — constructed bill-content report → PF_INFO.

Counterexamples:

Statistics deployed to condemn opponents' record → PF_ATTACK / PF_ADVOCACY, not neutral info.

«Comienza la votación. Efectuada la votación, dio el siguiente resultado: votos emitidos, 329; a favor, 197...» — proceeding-fact / vote management → PF_PROCEDURAL, not PF_INFO (§4.2).

Presidencia naming the next speaker or punto del orden del día → PF_PROCEDURAL, not PF_INFO.

Confused with	Resolution
PF_ADVOCACY	Info reports subject-matter facts; advocacy argues what should follow
PF_PROCEDURAL	See §4.2 — chamber/proceeding management vs policy/bill information
PF_DEFENSE	Correcting facts to rebut recoverable criticism → often DEFENSE (§4.3.4)

5.7 PF_DEFLECT — Deflection / evasion

Definition: Avoids question or topic; pivots without substantive engagement.

Decision rule: Speaker **does not address** the pending question/bill/challenge and shifts topic. **Rare** in prepared plenary speech; expect low frequency and higher adjudication. See §4.3.5.

Example (positive):

Pivot from bill substance to criticising an opponent's newspaper terminology instead of engaging the law.

— Patricia Reyes Rivera (Cs). Unit spdb-v1-unassigned-e022d9750540.

Illustrative constructed example (not corpus data):

Q (prior unit): «¿Apoyarán ustedes el artículo 2?» A: «Lo importante es hablar del caso de corrupción del otro partido.» — unanswered pivot → PF_DEFLECT.

Counterexample:

Direct rebuttal «No es así, señora Arrimadas» → PF_ATTACK / PF_DEFENSE, not deflect.

Confused with	Resolution
PF_ATTACK	See §4.3.5 — engage vs change subject
PF_ADVOCACY	Advocacy may criticise a bill; deflect avoids the pending question
PF_DEFENSE	See §4.3.5 — address vs sidestep

5.8 PF_PROCEDURAL — Ritual / procedural

Definition: Chamber management of the sitting: order of business, turn-taking, votes, timekeeping, suspension/resumption, agenda transitions, and formal ritual that organises debate — including factual statements **about the proceeding itself**.

Decision rule: Apply §4.2. Function of the utterance decides; chair identity does not. Formulaic courtesy alone inside a policy speech does **not** make the unit procedural (§4.1.2).

Examples (positive) — Wave-1 dual-agreement units (both annotators labelled PF_PROCEDURAL; not adjudicated gold):

«Comenzamos ahora, señorías, con el punto del orden del día [...] tiene la palabra [...] el señor Rodríguez Rodríguez.»

— Presidencia. Unit spdb-v1-unassigned-88da342f1158.

«Queda aprobado. Punto 10. Comienza la votación. Efectuada la votación, dio el siguiente resultado: votos emitidos, 329; a favor, 197; en contra, 1; abstenciones, 131.»

— Presidencia. Unit spdb-v1-unassigned-a462b78e5554. (Vote management + tally = proceeding-facts → PROCEDURAL, not INFO.)

«Muchas gracias. En el turno de fijación de posiciones, por el Grupo Parlamentario Mixto tiene la palabra, en primer lugar, el señor Martínez Oblanca.»

— Presidencia. Unit spdb-v1-unassigned-57ba0542f706.

Illustrative constructed example (not corpus data):

«Se suspende la sesión durante quince minutos. Se reanuda la sesión.» — suspension/resumption → PF_PROCEDURAL.

Counterexamples:

«Gracias, presidenta. [...] había 4,5 millones de españoles con graves problemas [...]» — courtesy opener + substance → label **substance**, not procedural.

A chair reading bill article content as exposition of policy substance (dominant move) → PF_INFO / other non-procedural label per tree — **not** automatic PROCEDURAL.

Confused with	Resolution
PF_INFO	See §4.2 — proceeding management vs subject-matter information
PF_PROPOSAL	Chair “vamos a votar” → PROCEDURAL; legislator pledge → PROPOSAL (§4.3.3)
PF_APPEAL	Ritual address to “señorías” about agenda → PROCEDURAL; appeal to citizens → APPEAL

6. Master decision tree — fallacies

Fallacies are **multi-label** (max **3**). Label only moves **explicit in the span** (±1 adjacent unit for context). Do not infer hidden intent.

START: Read unit (+ optional ±1 context).

1. Is there NO reasoning move matching any FAL_* definition below?
YES → FAL_NONE (fallacy_none_explicit=true; empty fallacy_labels)
NO → continue

2. For each suspected fallacy, verify definition in §7.
 - Ad hominem present? → FAL_ADHOM
 - Straw-man distortion present? → FAL_STRAW
 - False dilemma present? → FAL_DILEMMA
 - Slippery slope present? → FAL_SLOPE
 - Emotion/fear replaces evidence as primary warrant? → FAL_EMOTION
 - Hasty generalization present? → FAL_GENERAL
 - Whataboutism/tu quoque present? → FAL_WHATABOUT
3. If >3 fallacies suspected → keep top 3 by salience in this span.
4. Never combine FAL_NONE with any FAL_* label.
5. PF label is independent (code separately).

Evidence standard: A minimally competent reader of Spanish political discourse would agree the fallacy is present **in the span**.

Sarcasm/irony: Code the **literal rhetorical move**; flag `irony_possible=true` (analysis only, not a training label).

7. Political fallacy labels

Export: pipe-separated sorted IDs (e.g. FAL_ADHOM|FAL_WHATABOUT) or empty list + `fallacy_none_explicit=true`.

7.1 FAL_ADHOM — Ad hominem

Definition: Rejects or challenges a claim primarily via **personal attack** on the actor rather than the claim's merits.

Example (positive):

«Ustedes ni siquiera han estado sentados en la mesa donde se tomaban estas decisiones» — challenges Arrimadas' **standing**, not only the VAT policy merits.

Counterexample:

Evidence-based critique of a minister's **policy outcomes** with figures, without personal disqualification → no FAL_ADHOM (may still be PF_ATTACK).

Note: PF_ATTACK + FAL_ADHOM is common; PF_ATTACK without fallacy is valid when critique stays on policy/record with adequate evidential basis.

7.2 FAL_STRAW — Straw man / distortion

Definition: Misrepresents an opponent's position (stronger or weaker), then refutes the **distortion**.

Example (positive):

«Ustedes quieren eliminar todo el sector privado de la sanidad» when opponents proposed regulating one segment only — refuting an exaggerated caricature.

Counterexample:

Fair summary of opponent's published bill text followed by disagreement → not straw man.

7.3 FAL_DILEMMA — False dilemma

Definition: Presents two (or few) **exclusive** options when reasonable alternatives exist.

Example (positive):

«O están con la democracia española o están con el golpismo» — excludes intermediate positions or partial agreements.

Counterexample:

Binary vote on a **genuinely** two-option chamber division (sí/no to a single decree) described procedurally → not a false dilemma in argument.

7.4 FAL_SLOPE — Slippery slope

Definition: Claims an action **inevitably** leads to an extreme outcome without sufficient evidential steps.

Example (positive):

«Si permitimos esta enmienda, acabaremos disolviendo el Estado autonómico por completo» — no intermediate steps demonstrated.

Counterexample:

Cited chain of **documented** legislative precedents with explicit causal reasoning → not slippery slope (may be advocacy).

7.5 FAL_EMOTION — Appeal to emotion / fear

Definition: Replaces or evades evidentiary support with **fear, outrage, or pity** as the primary warrant.

Example (positive):

«Nuestros mayores morirán solos y abandonados» as the **main** reason to reject a bill, without policy engagement with the bill's measures.

Counterexample:

Citing **verified mortality statistics** from poverty-energy studies as evidence in policy advocacy → not emotion fallacy (may be PF_INFO / PF_ADVOCACY).

7.6 FAL_GENERAL — Hasty generalization

Definition: Generalizes from **unrepresentative or insufficient** evidence.

Example (positive):

«Un municipio gobernado por X colapsó; por tanto, X no puede gobernar España» — single case → national conclusion.

Counterexample:

Aggregate national statistics from official sources applied to a general claim → not hasty generalization.

7.7 FAL_WHATABOUT — Whataboutism / tu quoque

Definition: Deflects criticism by pointing to opponent's **unrelated or symmetric** behaviour without addressing the merit of the original charge.

Example (positive):

When asked about party financing transparency: «Y ustedes, ¿qué dicen del caso Gürtel?» — pivot without answering the question.

Counterexample:

Direct rebuttal with **on-topic** counter-evidence addressing the same policy domain → PF_DEFENSE, not whataboutism.

7.8 FAL_NONE — No fallacy

Definition: No fallacy from §7.1–7.7 is present in the span.

Export: fallacy_labels empty; fallacy_none_explicit=true. **Never** combine with any FAL_*.

8. Experimental labels (subset only)

Skip when experimental_subset=false.

ID	Name	Rule
SV_0	Semantic vacuity absent	Specifiable propositional political content
SV_1	Semantic vacuity present	Slogans, applause lines, content-free assent
SV_UNCLEAR	Unclear	< 3% of SV annotations
CA_0	Conceptual anachronism absent	Terms used in appropriate historical frame
CA_1	Conceptual anachronism present	Anachronistic projection onto past actors/events
CA_UNCLEAR	Unclear	< 3% of CA annotations

Experimental labels do **not** affect IberLEF leaderboard v1 ranking.

9. Edge cases compendium

#	Situation	PF guidance	Fallacy guidance
E1	Courtesy opener + long policy speech	Label substance	Code fallacies in full unit
E2	Chair turn with proceeding management	PF_PROCEDURAL if sitting-management dominates (§4.2); identity alone never decides	Usually FAL_NONE
E3	Quote of opponent then rebuttal	Label speaker's move , not quoted voice	Straw man if speaker distorts quote
E4	Rhetorical question attacking opponent	PF_ATTACK	FAL_EMOTION only if fear/outrage is primary warrant
E5	Manifesto programmatic list	Often PF_ADVOCACY or PF_PROPOSAL	Rare fallacies; avoid over-tagging
E6	Social-media post, sarcasm suspected	Code literal move; irony_possible=true	Do not guess hidden intent
E7	Policy attack with strong evidence	PF_ATTACK	No FAL_ADHOM if attack stays on merits/record
E8	«Todos los partidos menos nosotros»	Often PF_ATTACK or PF_APPEAL	Possible FAL_DILEMMA if exclusive framing
E9	Voting intention without new commitment	PF_ADVOCACY unless explicit pledge → PF_PROPOSAL	Usually none
E10	Multi-function long unit (segmentation error)	Pick dominant function; flag borderline	Top 3 fallacies by salience
E11	Speaker cites statistics from opponents	If correcting recoverable criticism → PF_DEFENSE (§4.3.4); if endorsing → PF_INFO/ PF_ADVOCACY	Rare
E12	Empty or near-empty span	Should not appear in pool; escalate to PI	—
E13	Vote tally / “queda aprobado”	Chamber management → usually PF_PROCEDURAL (§4.2), not PF_INFO	Usually FAL_NONE
E14	Mixed procedural shell + bill exposition	Substance if substance dominates; else PROCEDURAL; borderline if tied (§4.2)	Code fallacies on substance

10. Adjudication rules

10.1 Triggers (mandatory queue)

- Any **pragmatic function** disagreement between annotators A and B.
- Any **fallacy** disagreement when PF already agrees (or after PF adjudicated).
- Any unit with `borderline=true` on PF.
- Gold-item failures below QC thresholds (§11).

10.2 Procedure

1. Read adjudication_context (prior labels) and ± 1 context units.
2. Re-apply decision trees (§4, §6) **without** defaulting to majority vote.
3. Set gold labels using the same controls as primary annotation.
4. Fill **Adjudication notes** with brief rationale, especially when overriding both annotators.
5. Record gold in export: adjudicated=true; log label delta in adjudication_log.jsonl.
6. Complete within **72 hours** of queue assignment.

10.3 Precedence

Conflict type	Rule
PF split between advocacy vs attack	Apply §4 tree; check primary target (policy vs actor)
PF split defense vs attack	Reactive → defense; offensive → attack
PF procedural vs info	Apply §4.2 function test (not speaker identity); tallies/agenda → PROCEDURAL; bill/policy facts → INFO
Fallacy present vs none	Apply evidence standard ; when genuinely indeterminate after context, prefer none
>3 fallacies	Adjudicator selects top 3 by salience

10.4 Ontology changes

Adjudicators **must not** invent labels outside labels/*.tsv. Repeated confusion between the same pair (>10% of PF disagreements) triggers ontology/guideline revision (pilot go/no-go rule).

11. Quality control and reliability

11.1 Targets

Task	Metric	Codebook target	MVP release floor
Pragmatic function	Cohen's κ	> 0.80	≥ 0.67
Pragmatic function	Krippendorff's α	> 0.80	≥ 0.70
Fallacy (multi-label)	Macro-F1 (pairs)	≥ 0.60	≥ 0.48
Fallacy presence (binary)	Cohen's κ	≥ 0.70	≥ 0.52

11.2 Training protocol (to reach $\kappa > 0.80$)

1. Read this codebook and annotation/labelstudio/instructions.md.
2. Complete **5 practice units** (excluded from κ) with group review.
3. Annotate pilot batch; review disagreements with lead annotator.
4. Do **not** begin production until pilot PF $\kappa \geq 0.67$; recalibrate if below.
5. Monthly refresher on top confusion pairs from disagreement reports.

11.3 Hidden gold items

Maintain $\geq 85\%$ agreement on PF and $\geq 75\%$ on fallacies on 5% hidden gold subset during production.

11.4 Borderline budget

borderline=true on PF expected $\leq 5\%$ of submissions; higher rates trigger guideline clarification.

12. Quick reference — label IDs

Pragmatic function

ID	Name
PF_ADVOCACY	Policy advocacy
PF_ATTACK	Position attack
PF_DEFENSE	Defense / rebuttal
PF_PROPOSAL	Proposal / commitment
PF_APPEAL	Mobilization appeal
PF_INFO	Informational
PF_DEFLECT	Deflection / evasion
PF_PROCEDURAL	Ritual / procedural

Fallacies

ID	Name
FAL_ADHOM	Ad hominem
FAL_STRAW	Straw man / distortion
FAL_DILEMMA	False dilemma
FAL_SLOPE	Slippery slope
FAL_EMOTION	Appeal to emotion / fear
FAL_GENERAL	Hasty generalization
FAL_WHATABOUT	Whataboutism / tu quoque
FAL_NONE	No fallacy (explicit)

13. Related files

Path	Content
labels/pragmatic_functions.tsv	PF ontology (machine-readable; keep sync with §5)
labels/fallacies.tsv	Fallacy ontology
annotation/codebook/SPDB_Codebook_v1.md	

Path	Content
	Frozen Wave-1 instrument (codebook-v1.0.0)
annotation/guidelines/pragmatic_function_examples.md	Extended ParlaMint worked examples (guidelines-v1.1)
annotation/labelstudio/instructions.md	Label Studio UI workflow
annotation/attestation/annotator_session_attestation.md	Mandatory session provenance template
annotation/pilot_001/pilot_protocol.md	Wave-1 diagnostic pilot workflow (historical)
scripts/analysis/compute_cohen_kappa.py	Agreement analysis toolkit

14. Provenance / attestation (v1.1)

Before any reliability-set coding under this codebook, each annotator must complete annotation/attestation/annotator_session_attestation.md (or JSON twin). Past Wave-1 sessions remain **unknown / not recorded** — do not backfill attestation for Wave 1.

Wave-1 annotations remain v1.0-era diagnostic annotations and are not reclassified under v1.1.

Revision history

Version	Date	Changes
1.0.0	2026-06-18	Initial complete codebook: PF + fallacies, decision trees, adjudication, IA target $\kappa > 0.80$
1.1.0	2026-09-29	Candidate instrument for next reliability wave: PROCEDURAL↔INFO operational redefine (§4.2); high-value pair discriminators (§4.3); context policy (§2.1); example expansions; attestation pointer. Eight-class inventory unchanged. Wave-1 labels not reclassified.